

# BDA Lab

## Assignment-4

**Q Write a map reduce program to count the frequency of given word in a given file.**

Step 1 Start Hadoop Services

```
vboxuser@Ubuntu-hadoop:~$ start-dfs.sh
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [Ubuntu-hadoop]
vboxuser@Ubuntu-hadoop:~$ start-yarn.sh
Starting resourcemanager
Starting nodemanagers
vboxuser@Ubuntu-hadoop:~$ █
```

Step 2 create file

```
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab$ mkdir StudentAvg1
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab$ cd StudentAvg1
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab/StudentAvg$ nano StudentAverage.java
```

Step 3 Write word counting program

```
import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapreduce.*;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
```

```

public class StudentAverage {

    // Mapper Class
    public static class AvgMapper
        extends Mapper<LongWritable, Text, Text, FloatWritable> {

        private Text studentId = new Text();
        private FloatWritable grade = new FloatWritable();

        public void map(LongWritable key, Text value, Context context)
            throws IOException, InterruptedException {
            String line = value.toString();
            String[] parts = line.split(",");
            if (parts.length == 3) {
                studentId.set(parts[0]);
                grade.set(Float.parseFloat(parts[2]));
                context.write(studentId, grade);
            }
        }
    }

    // Reducer Class
    public static class AvgReducer
        extends Reducer<Text, FloatWritable, Text, FloatWritable> {

        private FloatWritable result = new FloatWritable();

        public void reduce(Text key, Iterable<FloatWritable> values,
                          Context context)
            throws IOException, InterruptedException {
            float sum = 0;

```

```

int count = 0;
for (FloatWritable val : values) {
    sum += val.get();
    count++;
}
float average = sum / count;
result.set(average);
context.write(key, result);
}

}

// Driver Class

public static void main(String[] args) throws Exception {
    Configuration conf = new Configuration();
    Job job = Job.getInstance(conf, "Average Grade");
    job.setJarByClass(StudentAverage.class);
    job.setMapperClass(AvgMapper.class);
    job.setReducerClass(AvgReducer.class);

    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(FloatWritable.class);

    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));

    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

## Step 4 Compile Program and create Jar

```
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab/StudentAvg$ javac -classpath `hadoop classpath` -d StudentAvg_classes StudentAverage.java
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab/StudentAvg$ jar -cvf studentavg.jar -C StudentAvg_classes/ .
added manifest
adding: StudentAverage.class(in = 1479) (out= 806)(deflated 45%)
adding: StudentAverage$AvgReducer.class(in = 1805) (out= 776)(deflated 57%)
adding: StudentAverage$AvgMapper.class(in = 1948) (out= 798)(deflated 59%)
```

## Step 5 Create Input file and upload to hdfs

```
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab/StudentAvg$ hdfs dfs -mkdir /studentinput
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab/StudentAvg$ hdfs dfs -put input.txt /studentinput
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab/StudentAvg$ hdfs dfs -ls /studentinput
Found 1 items
-rw-r--r--    3 vboxuser supergroup          66 2026-02-16 10:36 /studentinput/input.txt
```

101,Math,85

101,Science,90

102,Math,78

102,Science,82

103,Math,88

## Step 6 Run MapReduce Job

```
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab/StudentAvg$ hadoop jar studentavg.jar StudentAverage /studentinput /studentoutput
```

```
2026-02-16 10:39:06,173 INFO mapreduce.Job: map 0% reduce 0%
2026-02-16 10:40:02,561 INFO mapreduce.Job: map 100% reduce 0%
2026-02-16 10:40:39,011 INFO mapreduce.Job: map 100% reduce 100%
2026-02-16 10:41:09,548 INFO mapreduce.Job: Job job_1771237549532_0001 completed successfully
```

## Step 7 View Output

```
vboxuser@Ubuntu-hadoop:~/Desktop/BDA Lab/StudentAvg$ hdfs dfs -cat /studentoutput/part-r-00000
101      87.5
102      80.0
103      88.0
```